

```

1 import numpy as np
2
3 class Softmax(object):
4
5     def __init__(self, dims=[10, 3073]):
6         self.init_weights(dims=dims)
7
8     def init_weights(self, dims):
9         """
10        Initializes the weight matrix of the Softmax classifier.
11        Note that it has shape (C, D) where C is the number of
12        classes and D is the feature size.
13        """
14        self.W = np.random.normal(size=dims) * 0.0001
15
16    def loss(self, X, y):
17        """
18        Calculates the softmax loss.
19
20        Inputs have dimension D, there are C classes, and we operate on minibatches
21        of N examples.
22
23        Inputs:
24        - X: A numpy array of shape (N, D) containing a minibatch of data.
25        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
26            that X[i] has label c, where 0 ≤ c < C.
27
28        Returns a tuple of:
29        - loss as single float
30        """
31
32        # Initialize the loss to zero.
33        loss = 0.0
34
35        # ===== #
36        # YOUR CODE HERE:
37        # Calculate the normalized softmax loss. Store it as the variable loss.
38        # (That is, calculate the sum of the losses of all the training
39        # set margins, and then normalize the loss by the number of
40        # training examples.)
41        # ===== #
42
43        num_samples = X.shape[0]
44        num_classes = self.W.shape[0]
45
46        for i in np.arange(num_samples):
47
48            temp_log = 0
49            for j in np.arange(num_classes):
50                temp_log += np.exp(self.W[j].dot(X[i]))
51
52            loss += np.log(temp_log) - self.W[y[i]].dot(X[i])
53
54        loss = loss/num_samples
55
56        # ===== #
57        # END YOUR CODE HERE
58        # ===== #
59
60
61        return loss
62
63    def loss_and_grad(self, X, y):
64        """
65        Same as self.loss(X, y), except that it also returns the gradient.
66
67        Output: grad -- a matrix of the same dimensions as W containing
68            the gradient of the loss with respect to W.
69

```

```

70  """
71
72  # Initialize the loss and gradient to zero.
73  loss = 0.0
74  grad = np.zeros_like(self.W)
75
76  # ===== #
77  # YOUR CODE HERE:
78  # Calculate the softmax loss and the gradient. Store the gradient
79  # as the variable grad.
80  # ===== #
81
82  num_samples = X.shape[0]
83  num_classes = self.W.shape[0]
84
85  for i in np.arange(num_samples):
86
87      temp_log = 0
88      temp_grad = np.zeros_like(self.W[:,0])
89      for j in np.arange(num_classes):
90          score = np.exp(self.W[j].dot(X[i]))
91          temp_log += score
92          temp_grad[j] = score
93
94      temp_grad /= temp_log
95      temp_grad[y[i]] -= 1
96      grad += temp_grad[:, np.newaxis]*X[i]
97
98
99      loss += np.log(temp_log) - self.W[y[i]].dot(X[i])
100
101  loss = loss/num_samples
102  grad = grad/num_samples
103
104  # ===== #
105  # END YOUR CODE HERE
106  # ===== #
107
108  return loss, grad
109
110
111 def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
112  """
113  sample a few random elements and only return numerical
114  in these dimensions.
115  """
116
117  for i in np.arange(num_checks):
118      ix = tuple([np.random.randint(m) for m in self.W.shape])
119
120      oldval = self.W[ix]
121      self.W[ix] = oldval + h # increment by h
122      fxph = self.loss(X, y)
123      self.W[ix] = oldval - h # decrement by h
124      fxmh = self.loss(X,y) # evaluate f(x - h)
125      self.W[ix] = oldval # reset
126
127      grad_numerical = (fxph - fxmh) / (2 * h)
128      grad_analytic = your_grad[ix]
129      rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
130      print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))
131
132  def fast_loss_and_grad(self, X, y):
133  """
134  A vectorized implementation of loss_and_grad. It shares the same
135  inputs and outputs as loss_and_grad.
136  """
137  loss = 0.0
138  grad = np.zeros(self.W.shape) # initialize the gradient as zero
139

```

```

140 # ===== #
141 # YOUR CODE HERE:
142 # Calculate the softmax loss and gradient WITHOUT any for loops.
143 # ===== #
144
145 num_samples = X.shape[0]
146 num_classes = self.W.shape[0]
147
148 # Loss
149 scores = self.W.dot(X.T)
150 e_scores = np.exp(self.W.dot(X.T))
151 sums = np.sum(e_scores, axis=0)
152 log_sums = np.log(sums)
153 y_terms = scores[y, np.arange(num_samples)]
154 loss = np.sum(log_sums - y_terms)/num_samples
155
156
157 # Grad
158 e_scores = e_scores/sums
159 e_scores[y,np.arange(num_samples)] -= 1
160 grad = e_scores.dot(X)/num_samples
161
162 # ===== #
163 # END YOUR CODE HERE
164 # ===== #
165
166 return loss, grad
167
168 def train(self, X, y, learning_rate=1e-3, num_iters=100,
169         batch_size=200, verbose=False):
170     """
171     Train this linear classifier using stochastic gradient descent.
172
173     Inputs:
174     - X: A numpy array of shape (N, D) containing training data; there are N
175         training samples each of dimension D.
176     - y: A numpy array of shape (N,) containing training labels; y[i] = c
177         means that X[i] has label 0 ≤ c < C for C classes.
178     - learning_rate: (float) learning rate for optimization.
179     - num_iters: (integer) number of steps to take when optimizing
180     - batch_size: (integer) number of training examples to use at each step.
181     - verbose: (boolean) If true, print progress during optimization.
182
183     Outputs:
184     A list containing the value of the loss function at each training iteration.
185     """
186     num_train, dim = X.shape
187     num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
188
189     self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W
190
191     # Run stochastic gradient descent to optimize W
192     loss_history = []
193
194     for it in np.arange(num_iters):
195         X_batch = None
196         y_batch = None
197
198         # ===== #
199         # YOUR CODE HERE:
200         # Sample batch_size elements from the training data for use in
201         # gradient descent. After sampling,
202         # - X_batch should have shape: (dim, batch_size)
203         # - y_batch should have shape: (batch_size,)
204         # The indices should be randomly generated to reduce correlations
205         # in the dataset. Use np.random.choice. It's okay to sample with
206         # replacement.
207         # ===== #
208         mask = np.random.choice(np.arange(X.shape[0]), batch_size)
209         X_batch = X[mask]

```

```

210     y_batch = y[mask]
211     # ===== #
212     # END YOUR CODE HERE
213     # ===== #
214
215     # evaluate loss and gradient
216     loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
217     loss_history.append(loss)
218
219     # ===== #
220     # YOUR CODE HERE:
221     #   Update the parameters, self.W, with a gradient step
222     # ===== #
223     self.W -= learning_rate*grad
224
225     # ===== #
226     # END YOUR CODE HERE
227     # ===== #
228
229     if verbose and it % 100 == 0:
230         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
231
232     return loss_history
233
234 def predict(self, X):
235     """
236     Inputs:
237     - X: N x D array of training data. Each row is a D-dimensional point.
238
239     Returns:
240     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
241       array of length N, and each element is an integer giving the predicted
242       class.
243     """
244     y_pred = np.zeros(X.shape[1])
245     # ===== #
246     # YOUR CODE HERE:
247     #   Predict the labels given the training data.
248     # ===== #
249
250     scores = self.W.dot(X.T)
251     y_pred = np.argmax(scores,axis=0)
252
253
254
255     # ===== #
256     # END YOUR CODE HERE
257     # ===== #
258
259     return y_pred

```