

```

1 import numpy as np
2 import pdb
3
4 """
5 This code was based off of code from cs231n at Stanford University, and modified for ece239as at UCLA.
6 """
7 class SVM(object):
8
9     def __init__(self, dims=[10, 3073]):
10         self.init_weights(dims=dims)
11
12     def init_weights(self, dims):
13         """
14         Initializes the weight matrix of the SVM. Note that it has shape (C, D)
15         where C is the number of classes and D is the feature size.
16         """
17         self.W = np.random.normal(size=dims)
18
19     def loss(self, X, y):
20         """
21         Calculates the SVM loss.
22
23         Inputs have dimension D, there are C classes, and we operate on minibatches
24         of N examples.
25
26         Inputs:
27         - X: A numpy array of shape (N, D) containing a minibatch of data.
28         - y: A numpy array of shape (N,) containing training labels; y[i] = c means
29             that X[i] has label c, where 0 ≤ c < C.
30
31         Returns a tuple of:
32         - loss as single float
33         """
34
35         # compute the loss and the gradient
36         num_classes = self.W.shape[0]
37         num_train = X.shape[0]
38         loss = 0.0
39
40         for i in np.arange(num_train):
41
42             # ===== #
43             # YOUR CODE HERE:
44             # Calculate the normalized SVM loss, and store it as 'loss'.
45             # (That is, calculate the sum of the losses of all the training
46             # set margins, and then normalize the loss by the number of
47             # training examples.)
48             # ===== #
49
50             temp_loss = 0
51             for j in np.arange(num_classes):
52                 if j != y[i]:
53                     temp_loss += max(0, 1 + self.W[j].dot(X[i]) - self.W[y[i]].dot(X[i]))
54
55             loss += temp_loss
56
57         loss = loss/num_train
58
59         # ===== #
60         # END YOUR CODE HERE
61         # ===== #
62
63         return loss
64
65     def loss_and_grad(self, X, y):
66         """
67         Same as self.loss(X, y), except that it also returns the gradient.
68
69         Output: grad -- a matrix of the same dimensions as W containing

```

```

70     the gradient of the loss with respect to W.
71     """
72
73     # compute the loss and the gradient
74     num_classes = self.W.shape[0]
75     num_train = X.shape[0]
76     loss = 0.0
77     grad = np.zeros_like(self.W)
78
79     for i in np.arange(num_train):
80         # ===== #
81         # YOUR CODE HERE:
82         # Calculate the SVM loss and the gradient. Store the gradient in
83         # the variable grad.
84         # ===== #
85         temp_loss = 0
86
87         for j in np.arange(num_classes):
88             z = 1 + self.W[j].dot(X[i]) - self.W[y[i]].dot(X[i])
89
90             if j != y[i]:
91                 temp_loss += max(0, z)
92                 grad[j] += X[i] if z > 0 else 0
93                 grad[y[i]] += -1*X[i] if z > 0 else 0
94
95         loss += temp_loss
96
97
98
99         # ===== #
100        # END YOUR CODE HERE
101        # ===== #
102
103        loss /= num_train
104        grad /= num_train
105
106        return loss, grad
107
108    def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
109        """
110        sample a few random elements and only return numerical
111        in these dimensions.
112        """
113
114        for i in np.arange(num_checks):
115            ix = tuple([np.random.randint(m) for m in self.W.shape])
116
117            oldval = self.W[ix]
118            self.W[ix] = oldval + h # increment by h
119            fxph = self.loss(X, y)
120            self.W[ix] = oldval - h # decrement by h
121            fxmh = self.loss(X,y) # evaluate f(x - h)
122            self.W[ix] = oldval # reset
123
124            grad_numerical = (fxph - fxmh) / (2 * h)
125            grad_analytic = your_grad[ix]
126            rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
127            print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))
128
129    def fast_loss_and_grad(self, X, y):
130        """
131        A vectorized implementation of loss_and_grad. It shares the same
132        inputs and outputs as loss_and_grad.
133        """
134        loss = 0.0
135        grad = np.zeros(self.W.shape) # initialize the gradient as zero
136
137        # ===== #
138        # YOUR CODE HERE:
139        # Calculate the SVM loss WITHOUT any for loops.

```

```

140 # ===== #
141
142 loss_mat = self.W.dot(X.T)
143 loss_mat = loss_mat - loss_mat[y,np.arange(loss_mat.shape[1])] + 1
144 loss_mat[y, np.arange(loss_mat.shape[1])] = 0
145 loss_mat[loss_mat < 0] = 0
146 loss = loss_mat.sum()/X.shape[0]
147
148 # ===== #
149 # END YOUR CODE HERE
150 # ===== #
151
152
153 # ===== #
154 # YOUR CODE HERE:
155 # Calculate the SVM grad WITHOUT any for loops.
156 # ===== #
157 indicator = loss_mat
158 indicator[indicator>0] = 1
159 rsum = np.sum(indicator, axis=0)
160 indicator[y, np.arange(loss_mat.shape[1])] = -rsum
161 grad = indicator.dot(X)/X.shape[0]
162
163 # ===== #
164 # END YOUR CODE HERE
165 # ===== #
166
167 return loss, grad
168
169 def train(self, X, y, learning_rate=1e-3, num_iters=100,
170         batch_size=200, verbose=False):
171     """
172     Train this linear classifier using stochastic gradient descent.
173
174     Inputs:
175     - X: A numpy array of shape (N, D) containing training data; there are N
176         training samples each of dimension D.
177     - y: A numpy array of shape (N,) containing training labels; y[i] = c
178         means that X[i] has label 0 ≤ c < C for C classes.
179     - learning_rate: (float) learning rate for optimization.
180     - num_iters: (integer) number of steps to take when optimizing
181     - batch_size: (integer) number of training examples to use at each step.
182     - verbose: (boolean) If true, print progress during optimization.
183
184     Outputs:
185     A list containing the value of the loss function at each training iteration.
186     """
187     num_train, dim = X.shape
188     num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
189
190     self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W
191
192     # Run stochastic gradient descent to optimize W
193     loss_history = []
194
195     for it in np.arange(num_iters):
196         X_batch = None
197         y_batch = None
198
199         # ===== #
200         # YOUR CODE HERE:
201         # Sample batch_size elements from the training data for use in
202         # gradient descent. After sampling,
203         # - X_batch should have shape: (dim, batch_size)
204         # - y_batch should have shape: (batch_size,)
205         # The indices should be randomly generated to reduce correlations
206         # in the dataset. Use np.random.choice. It's okay to sample with
207         # replacement.
208         # ===== #
209         mask = np.random.choice(np.arange(X.shape[0]), batch_size)

```

```

210     X_batch = X[mask]
211     y_batch = y[mask]
212     # ===== #
213     # END YOUR CODE HERE
214     # ===== #
215
216     # evaluate loss and gradient
217     loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
218     loss_history.append(loss)
219
220     # ===== #
221     # YOUR CODE HERE:
222     #   Update the parameters, self.W, with a gradient step
223     # ===== #
224
225     self.W -= learning_rate*grad
226
227     # ===== #
228     # END YOUR CODE HERE
229     # ===== #
230
231     if verbose and it % 100 == 0:
232         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
233
234     return loss_history
235
236 def predict(self, X):
237     """
238     Inputs:
239     - X: N x D array of training data. Each row is a D-dimensional point.
240
241     Returns:
242     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
243       array of length N, and each element is an integer giving the predicted
244       class.
245     """
246     y_pred = np.zeros(X.shape[0])
247
248
249     # ===== #
250     # YOUR CODE HERE:
251     #   Predict the labels given the training data with the parameter self.W.
252     # ===== #
253
254     scores = self.W.dot(X.T)
255     y_pred = np.argmax(scores,axis=0)
256
257     # ===== #
258     # END YOUR CODE HERE
259     # ===== #
260
261     return y_pred

```