

```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6 """
7 This code was originally written for CS 231n at Stanford University
8 (cs231n.stanford.edu). It has been modified in various areas for use in the
9 ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
12 permission to use this code. To see the original version, please visit
13 cs231n.stanford.edu.
14 """
15
16 class TwoLayerNet(object):
17     """
18     A two-layer fully-connected neural network with ReLU nonlinearity and
19     softmax loss that uses a modular layer design. We assume an input dimension
20     of D, a hidden dimension of H, and perform classification over C classes.
21
22     The architecture should be affine - relu - affine - softmax.
23
24     Note that this class does not implement gradient descent; instead, it
25     will interact with a separate Solver object that is responsible for running
26     optimization.
27
28     The learnable parameters of the model are stored in the dictionary
29     self.params that maps parameter names to numpy arrays.
30     """
31
32     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
33                  dropout=0, weight_scale=1e-3, reg=0.0):
34         """
35         Initialize a new network.
36
37         Inputs:
38         - input_dim: An integer giving the size of the input
39         - hidden_dims: An integer giving the size of the hidden layer
40         - num_classes: An integer giving the number of classes to classify
41         - dropout: Scalar between 0 and 1 giving dropout strength.
42         - weight_scale: Scalar giving the standard deviation for random
43           initialization of the weights.
44         - reg: Scalar giving L2 regularization strength.
45         """
46         self.params = {}
47         self.reg = reg
48
49         # ===== #
50         # YOUR CODE HERE:
51         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
52         # self.params['W2'], self.params['b1'] and self.params['b2']. The
53         # biases are initialized to zero and the weights are initialized
54         # so that each parameter has mean 0 and standard deviation weight_scale.
55         # The dimensions of W1 should be (input_dim, hidden_dim) and the
56         # dimensions of W2 should be (hidden_dims, num_classes)
57         # ===== #
58
59         self.params['W1'] = weight_scale * np.random.randn(input_dim, hidden_dims)
60         self.params['b1'] = np.zeros(hidden_dims)
61         self.params['W2'] = weight_scale * np.random.randn(hidden_dims, num_classes)
62         self.params['b2'] = np.zeros(num_classes)
63
64         # ===== #
65         # END YOUR CODE HERE
66         # ===== #
67
68     def loss(self, X, y=None):
69         """
70         Compute loss and gradient for a minibatch of data.
71
72         Inputs:
73         - X: Array of input data of shape (N, d_1, ..., d_k)
74         - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
75
76         Returns:
77         If y is None, then run a test-time forward pass of the model and return:
78         - scores: Array of shape (N, C) giving classification scores, where
79           scores[i, c] is the classification score for X[i] and class c.
80
81         If y is not None, then run a training-time forward and backward pass and
82         return a tuple of:
83         - loss: Scalar value giving the loss
84         - grads: Dictionary with the same keys as self.params, mapping parameter
85           names to gradients of the loss with respect to those parameters.
86         """
87         scores = None
88
89         # ===== #
90         # YOUR CODE HERE:
91         # Implement the forward pass of the two-layer neural network. Store
92         # the class scores as the variable 'scores'. Be sure to use the layers
93         # you prior implemented.
94         # ===== #
95
96         h1, cache1 = affine_relu_forward(X, self.params['W1'], self.params['b1'])
97         scores, cache2 = affine_forward(h1, self.params['W2'], self.params['b2'])
98         # ===== #
99         # END YOUR CODE HERE
100         # ===== #
101
102         # If y is None then we are in test mode so just return scores
103         if y is None:

```

```

104     return scores
105
106     loss, grads = 0, {}
107     # ===== #
108     # YOUR CODE HERE:
109     # Implement the backward pass of the two-layer neural net. Store
110     # the loss as the variable 'loss' and store the gradients in the
111     # 'grads' dictionary. For the grads dictionary, grads['W1'] holds
112     # the gradient for W1, grads['b1'] holds the gradient for b1, etc.
113     # i.e., grads[k] holds the gradient for self.params[k].
114     #
115     # Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
116     # for each W. Be sure to include the 0.5 multiplying factor to
117     # match our implementation.
118     #
119     # And be sure to use the layers you prior implemented.
120     # ===== #
121     W1 = self.params['W1']
122     W2 = self.params['W2']
123
124     num_examples = scores.shape[0]
125
126     max_score = np.amax(scores, axis=1)
127     scores -= max_score[:, np.newaxis]
128
129     e_scores = np.exp(scores)
130     sums = np.sum(e_scores, axis=1)
131     log_sums = np.log(sums)
132     y_terms = scores[np.arange(num_examples), y]
133     loss = np.sum(log_sums - y_terms)/num_examples + .5*self.reg*np.sum(W1*W1) + .5*self.reg*np.sum(W2*W2)
134
135     d_scores = e_scores/sums[:, np.newaxis]
136     d_scores[np.arange(num_examples), y] -= 1
137     d_scores = d_scores.T/num_examples
138
139     dx2, dw2, db2 = affine_backward(d_scores.T, cache2)
140     dx1, dw1, db1 = affine_relu_backward(dx2, cache1)
141
142
143     grads['W1'] = dw1 + self.reg*W1
144     grads['b1'] = db1
145     grads['W2'] = dw2 + self.reg*W2
146     grads['b2'] = db2
147
148
149
150
151
152
153     # ===== #
154     # END YOUR CODE HERE
155     # ===== #
156
157     return loss, grads
158
159
160 class FullyConnectedNet(object):
161     """
162     A fully-connected neural network with an arbitrary number of hidden layers,
163     ReLU nonlinearities, and a softmax loss function. This will also implement
164     dropout and batch normalization as options. For a network with L layers,
165     the architecture will be
166
167     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
168
169     where batch normalization and dropout are optional, and the {...} block is
170     repeated L - 1 times.
171
172     Similar to the TwoLayerNet above, learnable parameters are stored in the
173     self.params dictionary and will be learned using the Solver class.
174     """
175
176     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
177                  dropout=0, use_batchnorm=False, reg=0.0,
178                  weight_scale=1e-2, dtype=np.float32, seed=None):
179         """
180         Initialize a new FullyConnectedNet.
181
182         Inputs:
183         - hidden_dims: A list of integers giving the size of each hidden layer.
184         - input_dim: An integer giving the size of the input.
185         - num_classes: An integer giving the number of classes to classify.
186         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
187           the network should not use dropout at all.
188         - use_batchnorm: Whether or not the network should use batch normalization.
189         - reg: Scalar giving L2 regularization strength.
190         - weight_scale: Scalar giving the standard deviation for random
191           initialization of the weights.
192         - dtype: A numpy datatype object; all computations will be performed using
193           this datatype. float32 is faster but less accurate, so you should use
194           float64 for numeric gradient checking.
195         - seed: If not None, then pass this random seed to the dropout layers. This
196           will make the dropout layers deterministic so we can gradient check the
197           model.
198         """
199         self.use_batchnorm = use_batchnorm
200         self.use_dropout = dropout > 0
201         self.reg = reg
202         self.num_layers = 1 + len(hidden_dims)
203         self.dtype = dtype
204         self.params = {}
205
206         # ===== #
207         # YOUR CODE HERE:

```

```

208 # Initialize all parameters of the network in the self.params dictionary.
209 # The weights and biases of layer 1 are W1 and b1; and in general the
210 # weights and biases of layer i are Wi and bi. The
211 # biases are initialized to zero and the weights are initialized
212 # so that each parameter has mean 0 and standard deviation weight_scale.
213 # ===== #
214
215 dimensions = [input_dim] + hidden_dims + [num_classes]
216
217 for i in np.arange(self.num_layers):
218     self.params['W{}'.format(i+1)] = weight_scale * np.random.randn(dimensions[i], dimensions[i+1])
219     self.params['b{}'.format(i+1)] = np.zeros(dimensions[i+1])
220
221
222
223 # ===== #
224 # END YOUR CODE HERE
225 # ===== #
226
227 # When using dropout we need to pass a dropout_param dictionary to each
228 # dropout layer so that the layer knows the dropout probability and the mode
229 # (train / test). You can pass the same dropout_param to each dropout layer.
230 self.dropout_param = {}
231 if self.use_dropout:
232     self.dropout_param = {'mode': 'train', 'p': dropout}
233     if seed is not None:
234         self.dropout_param['seed'] = seed
235
236 # With batch normalization we need to keep track of running means and
237 # variances, so we need to pass a special bn_param object to each batch
238 # normalization layer. You should pass self.bn_params[0] to the forward pass
239 # of the first batch normalization layer, self.bn_params[1] to the forward
240 # pass of the second batch normalization layer, etc.
241 self.bn_params = []
242 if self.use_batchnorm:
243     self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]
244
245 # Cast all parameters to the correct datatype
246 for k, v in self.params.items():
247     self.params[k] = v.astype(dtype)
248
249
250 def loss(self, X, y=None):
251     """
252     Compute loss and gradient for the fully-connected net.
253
254     Input / output: Same as TwoLayerNet above.
255     """
256     X = X.astype(self.dtype)
257     mode = 'test' if y is None else 'train'
258
259     # Set train/test mode for batchnorm params and dropout param since they
260     # behave differently during training and testing.
261     if self.dropout_param is not None:
262         self.dropout_param['mode'] = mode
263     if self.use_batchnorm:
264         for bn_param in self.bn_params:
265             bn_param[mode] = mode
266
267     scores = None
268
269     # ===== #
270     # YOUR CODE HERE:
271     # Implement the forward pass of the FC net and store the output
272     # scores as the variable "scores".
273     # ===== #
274     caches = []
275     layer_scores = []
276
277     layer_scores.append(X)
278
279     for i in np.arange(self.num_layers-1):
280         temp_score, temp_cache = affine_relu_forward(layer_scores[i], self.params['W{}'.format(i+1)], self.params['b{}'.format(i+1)])
281         caches.append(temp_cache)
282         layer_scores.append(temp_score)
283
284     temp_score, temp_cache = affine_forward(layer_scores[self.num_layers-1], self.params['W{}'.format(self.num_layers)], self.params['b{}'.format(self.num_layers)])
285     caches.append(temp_cache)
286     layer_scores.append(temp_score)
287
288     scores = layer_scores[-1]
289
290     # ===== #
291     # END YOUR CODE HERE
292     # ===== #
293
294     # If test mode return early
295     if mode == 'test':
296         return scores
297
298     loss, grads = 0.0, {}
299     # ===== #
300     # YOUR CODE HERE:
301     # Implement the backwards pass of the FC net and store the gradients
302     # in the grads dict, so that grads[k] is the gradient of self.params[k]
303     # Be sure your L2 regularization includes a 0.5 factor.
304     # ===== #
305     num_examples = scores.shape[0]
306
307     max_score = np.amax(scores, axis=1)
308     scores -= max_score[:, np.newaxis]
309
310     e_scores = np.exp(scores)
311     sums = np.sum(e_scores, axis=1)

```

```

312 log_sums = np.log(sums)
313 y_terms = scores[np.arange(num_examples), y]
314
315 reg_loss = 0
316 for i in np.arange(self.num_layers):
317     W = self.params['W{}'.format(i+1)]
318     reg_loss += .5*self.reg*np.sum(W*W)
319
320 loss = np.sum(log_sums - y_terms)/num_examples + reg_loss
321
322
323
324 d_scores = e_scores/sums[:,np.newaxis]
325 d_scores[np.arange(num_examples),y] -= 1
326 d_scores = d_scores/num_examples
327
328 #print len(caches)
329 #print self.num_layers
330
331 dx, dw, db = affine_backward(d_scores, caches[self.num_layers-1])
332 grads['W{}'.format(self.num_layers)] = dw + self.reg*self.params['W{}'.format(self.num_layers)]
333 grads['b{}'.format(self.num_layers)] = db
334 d_scores = dx
335
336
337 for i in np.arange(self.num_layers-2, -1,-1):
338     dx, dw, db = affine_relu_backward(d_scores, caches[i])
339     grads['W{}'.format(i+1)] = dw + self.reg*self.params['W{}'.format(i+1)]
340     grads['b{}'.format(i+1)] = db
341     d_scores = dx
342
343 #dx2, dw2, db2 = affine_relu_backward(d_scores, cache[i])
344 # dx1, dw1, db1 = affine_relu_backward(dx2, cache1)
345
346
347 # ===== #
348 # END YOUR CODE HERE
349 # ===== #
350 return loss, grads

```