

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
10 permission to use this code. To see the original version, please visit
11 cs231n.stanford.edu.
12 """
13
14 class TwoLayerNet(object):
15     """
16     A two-layer fully-connected neural network. The net has an input dimension of
17     N, a hidden layer dimension of H, and performs classification over C classes.
18     We train the network with a softmax loss function and L2 regularization on the
19     weight matrices. The network uses a ReLU nonlinearity after the first fully
20     connected layer.
21
22     In other words, the network has the following architecture:
23
24     input - fully connected layer - ReLU - fully connected layer - softmax
25
26     The outputs of the second fully-connected layer are the scores for each class.
27     """
28
29     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
30         """
31         Initialize the model. Weights are initialized to small random values and
32         biases are initialized to zero. Weights and biases are stored in the
33         variable self.params, which is a dictionary with the following keys:
34
35         W1: First layer weights; has shape (H, D)
36         b1: First layer biases; has shape (H,)
37         W2: Second layer weights; has shape (C, H)
38         b2: Second layer biases; has shape (C,)
39
40         Inputs:
41         - input_size: The dimension D of the input data.
42         - hidden_size: The number of neurons H in the hidden layer.
43         - output_size: The number of classes C.
44         """
45         self.params = {}
46         self.params['W1'] = std * np.random.randn(hidden_size, input_size)
47         self.params['b1'] = np.zeros(hidden_size)
48         self.params['W2'] = std * np.random.randn(output_size, hidden_size)
49         self.params['b2'] = np.zeros(output_size)
50
51     def loss(self, X, y=None, reg=0.0):
52         """
53         Compute the loss and gradients for a two layer fully connected neural
54         network.
55
56         Inputs:
57         - X: Input data of shape (N, D). Each X[i] is a training sample.
58         - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
59             an integer in the range 0 <= y[i] < C. This parameter is optional; if it
60             is not passed then we only return scores, and if it is passed then we
61             instead return the loss and gradients.
62         - reg: Regularization strength.
63
64         Returns:
65         If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
66             the score for class c on input X[i].
67
68         If y is not None, instead return a tuple of:

```

```

70 - loss: Loss (data loss and regularization loss) for this batch of training
71     samples.
72 - grads: Dictionary mapping parameter names to gradients of those parameters
73     with respect to the loss function; has the same keys as self.params.
74     """
75     # Unpack variables from the params dictionary
76     W1, b1 = self.params['W1'], self.params['b1']
77     W2, b2 = self.params['W2'], self.params['b2']
78     N, D = X.shape
79
80     # Compute the forward pass
81     scores = None
82
83     # ===== #
84     # YOUR CODE HERE:
85     # Calculate the output scores of the neural network. The result
86     # should be (C, N). As stated in the description for this class,
87     # there should not be a ReLU layer after the second FC layer.
88     # The output of the second FC layer is the output scores. Do not
89     # use a for loop in your implementation.
90     # ===== #
91
92     H1 = W1.dot(X.T) + b1[:, np.newaxis]
93     H1[H1<0] = 0
94     scores = W2.dot(H1) + b2[:, np.newaxis]
95     scores = scores.T
96
97     # ===== #
98     # END YOUR CODE HERE
99     # ===== #
100
101
102     # If the targets are not given then jump out, we're done
103     if y is None:
104         return scores
105
106     # Compute the loss
107     loss = None
108
109     # ===== #
110     # YOUR CODE HERE:
111     # Calculate the loss of the neural network. This includes the
112     # softmax loss and the L2 regularization for W1 and W2. Store the
113     # total loss in the variable loss. Multiply the regularization
114     # loss by 0.5 (in addition to the factor reg).
115     # ===== #
116
117     # scores is num_examples by num_classes
118     num_examples = scores.shape[0]
119
120     max_score = np.amax(scores, axis=1)
121     scores -= max_score[:, np.newaxis]
122
123     e_scores = np.exp(scores)
124     sums = np.sum(e_scores, axis=1)
125     log_sums = np.log(sums)
126     y_terms = scores[np.arange(num_examples), y]
127     loss = np.sum(log_sums - y_terms)/num_examples + .5*reg*np.sum(W1*W1) + .5*reg*np.sum(W2*W2)
128     # ===== #
129     # END YOUR CODE HERE
130     # ===== #
131
132     grads = {}
133
134     # ===== #
135     # YOUR CODE HERE:
136     # Implement the backward pass. Compute the derivatives of the
137     # weights and the biases. Store the results in the grads
138     # dictionary. e.g., grads['W1'] should store the gradient for

```

```

139 # W1, and be of the same size as W1.
140 # ===== #
141
142 #print W1.shape
143 #print H1.shape
144 #print scores.shape
145
146 d_scores = e_scores/sums[:,np.newaxis]
147 d_scores[np.arange(num_examples),y] -= 1
148 d_scores = d_scores.T/num_examples
149
150 b2_grad = np.sum(d_scores,axis=1)
151 W2_grad = d_scores.dot(H1.T)
152
153 r_grad = W2.T.dot(d_scores)
154 r_grad[H1<=0] = 0
155
156 b1_grad = np.sum(r_grad,axis=1)
157 W1_grad = r_grad.dot(X)
158
159 grads['b1'] = b1_grad
160 grads['W1'] = W1_grad + reg*W1
161
162 grads['b2'] = b2_grad
163 grads['W2'] = W2_grad + reg*W2
164
165 # ===== #
166 # END YOUR CODE HERE
167 # ===== #
168
169 return loss, grads
170
171 def train(self, X, y, X_val, y_val,
172         learning_rate=1e-3, learning_rate_decay=0.95,
173         reg=1e-5, num_iters=100,
174         batch_size=200, verbose=False):
175     """
176     Train this neural network using stochastic gradient descent.
177
178     Inputs:
179     - X: A numpy array of shape (N, D) giving training data.
180     - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
181         X[i] has label c, where 0 <= c < C.
182     - X_val: A numpy array of shape (N_val, D) giving validation data.
183     - y_val: A numpy array of shape (N_val,) giving validation labels.
184     - learning_rate: Scalar giving learning rate for optimization.
185     - learning_rate_decay: Scalar giving factor used to decay the learning rate
186         after each epoch.
187     - reg: Scalar giving regularization strength.
188     - num_iters: Number of steps to take when optimizing.
189     - batch_size: Number of training examples to use per step.
190     - verbose: boolean; if true print progress during optimization.
191     """
192     num_train = X.shape[0]
193     iterations_per_epoch = max(num_train / batch_size, 1)
194
195     # Use SGD to optimize the parameters in self.model
196     loss_history = []
197     train_acc_history = []
198     val_acc_history = []
199
200     for it in np.arange(num_iters):
201         X_batch = None
202         y_batch = None
203
204         # ===== #
205         # YOUR CODE HERE:
206         # Create a minibatch by sampling batch_size samples randomly.
207         # ===== #

```

```

208 mask = np.random.choice(np.arange(X.shape[0]), batch_size)
209 X_batch = X[mask]
210 y_batch = y[mask]
211
212 # ===== #
213 # END YOUR CODE HERE
214 # ===== #
215
216 # Compute loss and gradients using the current minibatch
217 loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
218 loss_history.append(loss)
219
220 # ===== #
221 # YOUR CODE HERE:
222 # Perform a gradient descent step using the minibatch to update
223 # all parameters (i.e., W1, W2, b1, and b2).
224 # ===== #
225
226 self.params['W1'] -= learning_rate*grads['W1']
227 self.params['b1'] -= learning_rate*grads['b1']
228 self.params['W2'] -= learning_rate*grads['W2']
229 self.params['b2'] -= learning_rate*grads['b2']
230
231
232 # ===== #
233 # END YOUR CODE HERE
234 # ===== #
235
236 if verbose and it % 100 == 0:
237     print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
238
239 # Every epoch, check train and val accuracy and decay learning rate.
240 if it % iterations_per_epoch == 0:
241     # Check accuracy
242     train_acc = (self.predict(X_batch) == y_batch).mean()
243     val_acc = (self.predict(X_val) == y_val).mean()
244     train_acc_history.append(train_acc)
245     val_acc_history.append(val_acc)
246
247     # Decay learning rate
248     learning_rate *= learning_rate_decay
249
250 return {
251     'loss_history': loss_history,
252     'train_acc_history': train_acc_history,
253     'val_acc_history': val_acc_history,
254 }
255
256 def predict(self, X):
257     """
258     Use the trained weights of this two-layer network to predict labels for
259     data points. For each data point we predict scores for each of the C
260     classes, and assign each data point to the class with the highest score.
261
262     Inputs:
263     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
264         classify.
265
266     Returns:
267     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
268         the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
269         to have class c, where 0 <= c < C.
270     """
271     y_pred = None
272
273     # ===== #
274     # YOUR CODE HERE:
275     # Predict the class given the input data.
276     # ===== #

```

```
277 W1, b1 = self.params['W1'], self.params['b1']
278 W2, b2 = self.params['W2'], self.params['b2']
279
280 N, D = X.shape
281 H1 = W1.dot(X.T) + b1[:, np.newaxis]
282 H1[H1<0] = 0
283 scores = W2.dot(H1) + b2[:, np.newaxis]
284
285 y_pred = np.argmax(scores, axis=0)
286
287
288 # ===== #
289 # END YOUR CODE HERE
290 # ===== #
291
292 return y_pred
```