

```

1 import numpy as np
2 import pdb
3
4 from .layers import *
5 from .layer_utils import *
6
7 """
8 This code was originally written for CS 231n at Stanford University
9 (cs231n.stanford.edu). It has been modified in various areas for use in the
10 ECE 239AS class at UCLA. This includes the descriptions of what code to
11 implement as well as some slight potential changes in variable names to be
12 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
13 permission to use this code. To see the original version, please visit
14 cs231n.stanford.edu.
15 """
16
17 class FullyConnectedNet(object):
18     """
19     A fully-connected neural network with an arbitrary number of hidden layers,
20     ReLU nonlinearities, and a softmax loss function. This will also implement
21     dropout and batch normalization as options. For a network with L layers,
22     the architecture will be
23
24     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
25
26     where batch normalization and dropout are optional, and the {...} block is
27     repeated L - 1 times.
28
29     Similar to the TwoLayerNet above, learnable parameters are stored in the
30     self.params dictionary and will be learned using the Solver class.
31     """
32
33     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
34                 dropout=0, use_batchnorm=False, reg=0.0,
35                 weight_scale=1e-2, dtype=np.float32, seed=None):
36         """
37         Initialize a new FullyConnectedNet.
38
39         Inputs:
40         - hidden_dims: A list of integers giving the size of each hidden layer.
41         - input_dim: An integer giving the size of the input.
42         - num_classes: An integer giving the number of classes to classify.
43         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
44           the network should not use dropout at all.
45         - use_batchnorm: Whether or not the network should use batch normalization.
46         - reg: Scalar giving L2 regularization strength.
47         - weight_scale: Scalar giving the standard deviation for random
48           initialization of the weights.
49         - dtype: A numpy datatype object; all computations will be performed using
50           this datatype. float32 is faster but less accurate, so you should use
51           float64 for numeric gradient checking.
52         - seed: If not None, then pass this random seed to the dropout layers. This
53           will make the dropout layers deterministic so we can gradient check the
54           model.
55         """
56         self.use_batchnorm = use_batchnorm
57         self.use_dropout = dropout > 0
58         self.reg = reg
59         self.num_layers = 1 + len(hidden_dims)
60         self.dtype = dtype
61         self.params = {}
62
63         # ===== #
64         # YOUR CODE HERE:
65         # Initialize all parameters of the network in the self.params dictionary.
66         # The weights and biases of layer 1 are W1 and b1; and in general the
67         # weights and biases of layer i are Wi and bi. The
68         # biases are initialized to zero and the weights are initialized
69         # so that each parameter has mean 0 and standard deviation weight_scale.
70         #
71         # BATCHNORM: Initialize the gammas of each layer to 1 and the beta
72         # parameters to zero. The gamma and beta parameters for layer 1 should
73         # be self.params['gamma1'] and self.params['beta1']. For layer 2, they
74         # should be gamma2 and beta2, etc. Only use batchnorm if self.use_batchnorm
75         # is true and DO NOT batch normalize the output scores.
76         # ===== #
77         dimensions = [input_dim] + hidden_dims + [num_classes]
78
79         for i in np.arange(self.num_layers):
80             self.params['W{}'.format(i+1)] = weight_scale * np.random.randn(dimensions[i], dimensions[i+1])
81             self.params['b{}'.format(i+1)] = np.zeros(dimensions[i+1])
82             if self.use_batchnorm and i < self.num_layers-1:
83                 self.params['gamma{}'.format(i+1)] = np.ones((1,dimensions[i+1]))
84                 self.params['beta{}'.format(i+1)] = np.zeros((1,dimensions[i+1]))
85
86
87         # ===== #
88         # END YOUR CODE HERE
89         # ===== #
90
91         # When using dropout we need to pass a dropout_param dictionary to each
92         # dropout layer so that the layer knows the dropout probability and the mode
93         # (train / test). You can pass the same dropout_param to each dropout layer.
94         self.dropout_param = {}
95         if self.use_dropout:
96             self.dropout_param = {'mode': 'train', 'p': dropout}
97             if seed is not None:
98                 self.dropout_param['seed'] = seed
99
100

```

```

101 # With batch normalization we need to keep track of running means and
102 # variances, so we need to pass a special bn_param object to each batch
103 # normalization layer. You should pass self.bn_params[0] to the forward pass
104 # of the first batch normalization layer, self.bn_params[1] to the forward
105 # pass of the second batch normalization layer, etc.
106 self.bn_params = []
107 if self.use_batchnorm:
108     self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]
109
110 # Cast all parameters to the correct datatype
111 for k, v in self.params.items():
112     self.params[k] = v.astype(dtype)
113
114
115 def loss(self, X, y=None):
116     """
117     Compute loss and gradient for the fully-connected net.
118
119     Input / output: Same as TwoLayerNet above.
120     """
121     X = X.astype(self.dtype)
122     mode = 'test' if y is None else 'train'
123
124     # Set train/test mode for batchnorm params and dropout param since they
125     # behave differently during training and testing.
126     if self.dropout_param is not None:
127         self.dropout_param['mode'] = mode
128     if self.use_batchnorm:
129         for bn_param in self.bn_params:
130             bn_param['mode'] = mode
131
132     scores = None
133
134     # ===== #
135     # YOUR CODE HERE:
136     # Implement the forward pass of the FC net and store the output
137     # scores as the variable "scores".
138     #
139     # BATCHNORM: If self.use_batchnorm is true, insert a batchnorm layer
140     # between the affine_forward and relu_forward layers. You may
141     # also write an affine_batchnorm_relu() function in layer_utils.py.
142     #
143     # DROPOUT: If dropout is non-zero, insert a dropout layer after
144     # every ReLU layer.
145     # ===== #
146     caches = []
147     layer_scores = []
148     do_caches = []
149     do_cache = None
150
151     layer_scores.append(X)
152     temp_score = X
153
154     for i in np.arange(self.num_layers-1):
155         if self.use_batchnorm:
156             temp_score, temp_cache = affine_batchnorm_relu(temp_score, self.params['W{}'.format(i+1)],
157                 self.params['b{}'.format(i+1)], self.params['gamma{}'.format(i+1)],
158                 self.params['beta{}'.format(i+1)], self.bn_params[i])
159         else:
160             temp_score, temp_cache = affine_relu_forward(temp_score, self.params['W{}'.format(i+1)],
161                 self.params['b{}'.format(i+1)])
162         caches.append(temp_cache)
163         layer_scores.append(temp_score)
164         if self.use_dropout:
165             temp_score, do_cache = dropout_forward(temp_score, self.dropout_param)
166             do_caches.append(do_cache)
167
168
169
170
171 temp_score, temp_cache = affine_forward(temp_score, self.params['W{}'.format(self.num_layers)], self.params['b{}'.format(self.num_layers)])
172 caches.append(temp_cache)
173 layer_scores.append(temp_score)
174
175 scores = layer_scores[-1]
176
177 # ===== #
178 # END YOUR CODE HERE
179 # ===== #
180
181 # If test mode return early
182 if mode == 'test':
183     return scores
184
185 loss, grads = 0.0, {}
186 # ===== #
187 # YOUR CODE HERE:
188 # Implement the backwards pass of the FC net and store the gradients
189 # in the grads dict, so that grads[k] is the gradient of self.params[k]
190 # Be sure your L2 regularization includes a 0.5 factor.
191 #
192 # BATCHNORM: Incorporate the backward pass of the batchnorm.
193 #
194 # DROPOUT: Incorporate the backward pass of dropout.
195 # ===== #
196 num_examples = scores.shape[0]
197
198 max_score = np.amax(scores, axis=1)
199 scores -= max_score[:, np.newaxis]
200
201 e_scores = np.exp(scores)

```

```

202     sums = np.sum(e_scores, axis=1)
203     log_sums = np.log(sums)
204     y_terms = scores[np.arange(num_examples), y]
205
206     reg_loss = 0
207     for i in np.arange(self.num_layers):
208         W = self.params['W{}'.format(i+1)]
209         reg_loss += .5*self.reg*np.sum(W*W)
210
211     loss = np.sum(log_sums - y_terms)/num_examples + reg_loss
212
213     d_scores = e_scores/sums[:,np.newaxis]
214     d_scores[np.arange(num_examples),y] -= 1
215     d_scores = d_scores/num_examples
216
217     dx, dw, db = affine_backward(d_scores, caches[self.num_layers-1])
218     grads['W{}'.format(self.num_layers)] = dw + self.reg*self.params['W{}'.format(self.num_layers)]
219     grads['b{}'.format(self.num_layers)] = db
220
221     for i in np.arange(self.num_layers-2, -1,-1):
222         if self.use_dropout:
223             dx = dropout_backward(dx, do_caches[i])
224         if self.use_batchnorm:
225             dx, dw, db, dgamma, dbeta = affine_batchnorm_relu_back(dx, caches[i])
226             grads['gamma{}'.format(i+1)] = dgamma
227             grads['beta{}'.format(i+1)] = dbeta
228         else:
229             dx, dw, db = affine_relu_backward(dx, caches[i])
230             grads['W{}'.format(i+1)] = dw + self.reg*self.params['W{}'.format(i+1)]
231             grads['b{}'.format(i+1)] = db
232
233
234
235     # ===== #
236     # END YOUR CODE HERE
237     # ===== #
238
239     return loss, grads

```