

```

1 import numpy as np
2
3 from nndl.layers import *
4 from nndl.conv_layers import *
5 from cs231n.fast_layers import *
6 from nndl.layer_utils import *
7 from nndl.conv_layer_utils import *
8
9 import pdb
10
11 """
12 This code was originally written for CS 231n at Stanford University
13 (cs231n.stanford.edu). It has been modified in various areas for use in the
14 ECE 239AS class at UCLA. This includes the descriptions of what code to
15 implement as well as some slight potential changes in variable names to be
16 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
17 permission to use this code. To see the original version, please visit
18 cs231n.stanford.edu.
19 """
20
21 class ThreeLayerConvNet(object):
22     """
23     A three-layer convolutional network with the following architecture:
24
25     conv - relu - 2x2 max pool - affine - relu - affine - softmax
26
27     The network operates on minibatches of data that have shape (N, C, H, W)
28     consisting of N images, each with height H and width W and with C input
29     channels.
30     """
31
32     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
33                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
34                 dtype=np.float32, use_batchnorm=False):
35         """
36         Initialize a new network.
37
38         Inputs:
39         - input_dim: Tuple (C, H, W) giving size of input data
40         - num_filters: Number of filters to use in the convolutional layer
41         - filter_size: Size of filters to use in the convolutional layer
42         - hidden_dim: Number of units to use in the fully-connected hidden layer
43         - num_classes: Number of scores to produce from the final affine layer.
44         - weight_scale: Scalar giving standard deviation for random initialization
45           of weights.
46         - reg: Scalar giving L2 regularization strength
47         - dtype: numpy datatype to use for computation.
48         """
49         self.use_batchnorm = use_batchnorm
50         self.params = {}
51         self.reg = reg
52         self.dtype = dtype
53
54         # ===== #
55         # YOUR CODE HERE:
56         # Initialize the weights and biases of a three layer CNN. To initialize:
57         # - the biases should be initialized to zeros.
58         # - the weights should be initialized to a matrix with entries
59         #   drawn from a Gaussian distribution with zero mean and
60         #   standard deviation given by weight_scale.
61         # ===== #
62
63         self.params['W1'] = weight_scale * np.random.randn(num_filters, input_dim[0], filter_size, filter_size)
64         self.params['b1'] = np.zeros(num_filters)
65         self.params['W2'] = weight_scale * np.random.randn(num_filters*(input_dim[1]/2)*(input_dim[2]/2), hidden_dim)
66         self.params['b2'] = np.zeros(hidden_dim)
67         self.params['W3'] = weight_scale * np.random.randn(hidden_dim, num_classes)
68         self.params['b3'] = np.zeros(num_classes)
69
70         # ===== #
71         # END YOUR CODE HERE
72         # ===== #
73
74         for k, v in self.params.items():
75             self.params[k] = v.astype(dtype)
76
77
78     def loss(self, X, y=None):

```

```

80     """
81     Evaluate loss and gradient for the three-layer convolutional network.
82
83     Input / output: Same API as TwoLayerNet in fc_net.py.
84     """
85     W1, b1 = self.params['W1'], self.params['b1']
86     W2, b2 = self.params['W2'], self.params['b2']
87     W3, b3 = self.params['W3'], self.params['b3']
88
89     # pass conv_param to the forward pass for the convolutional layer
90     filter_size = W1.shape[2]
91     conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
92
93     # pass pool_param to the forward pass for the max-pooling layer
94     pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
95
96     scores = None
97
98     # ===== #
99     # YOUR CODE HERE:
100    # Implement the forward pass of the three layer CNN. Store the output
101    # scores as the variable "scores".
102    # ===== #
103
104    c1, conv_cache1 = conv_forward_fast(X, W1, b1, conv_param)
105    h1, r_cache1 = relu_forward(c1)
106    mp1, mp_cache1 = max_pool_forward_fast(h1, pool_param)
107    h2, aff_cache1 = affine_relu_forward(mp1, W2, b2)
108    scores, aff_cache2 = affine_forward(h2, W3, b3)
109
110    # ===== #
111    # END YOUR CODE HERE
112    # ===== #
113
114    if y is None:
115        return scores
116
117    loss, grads = 0, {}
118    # ===== #
119    # YOUR CODE HERE:
120    # Implement the backward pass of the three layer CNN. Store the grads
121    # in the grads dictionary, exactly as before (i.e., the gradient of
122    # self.params[k] will be grads[k]). Store the loss as "loss", and
123    # don't forget to add regularization on ALL weight matrices.
124    # ===== #
125
126    loss, dout = softmax_loss(scores, y)
127    loss += .5*self.reg*np.sum(W1*W1) + .5*self.reg*np.sum(W2*W2) + .5*self.reg*np.sum(W3*W3)
128
129    dx3, dw3, db3 = affine_backward(dout, aff_cache2)
130    dx2, dw2, db2 = affine_relu_backward(dx3, aff_cache1)
131    dmp = max_pool_backward_fast(dx2, mp_cache1)
132    dr = relu_backward(dmp, r_cache1)
133    dx1, dw1, db1 = conv_backward_fast(dr, conv_cache1)
134
135
136    grads['W1'] = dw1 + self.reg*W1
137    grads['b1'] = db1
138    grads['W2'] = dw2 + self.reg*W2
139    grads['b2'] = db2
140    grads['W3'] = dw3 + self.reg*W3
141    grads['b3'] = db3
142
143
144    # ===== #
145    # END YOUR CODE HERE
146    # ===== #
147
148    return loss, grads
149
150

```