

Convolutional neural network layers

In this notebook, we will build the convolutional neural network layers. This will be followed by a spatial batchnorm, and then in the final notebook of this assignment, we will train a CNN to further improve the validation accuracy on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

```
In [1]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Implementing CNN layers

Just as we implemented modular layers for fully connected networks, batch normalization, and dropout, we'll want to implement modular layers for convolutional neural networks. These layers are in `nndl/conv_layers.py`.

Convolutional forward pass

Begin by implementing a naive version of the forward pass of the CNN that uses for loops. This function is `conv_forward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a triple for loop.

After you implement `conv_forward_naive`, test your implementation by running the cell below.

```
In [2]: x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [[ 0.21027089,  0.21661097],
                           [ 0.22847626,  0.23004637]],
                          [[ 0.50813986,  0.54309974],
                           [ 0.64082444,  0.67101435]]],
                         [[[-0.98053589, -1.03143541],
                           [-1.19128892, -1.24695841]],
                          [[ 0.69108355,  0.66880383],
                           [ 0.59480972,  0.56776003]],
                          [[ 2.36270298,  2.36904306],
                           [ 2.38090835,  2.38247847]]]])

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))

Testing conv_forward_naive
('difference: ', 2.2121476417505994e-08)
```

Convolutional backward pass

Now, implement a naive version of the backward pass of the CNN. The function is `conv_backward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a quadruple for loop.

After you implement `conv_backward_naive`, test your implementation by running the cell below.

```

In [3]: x = np.random.randn(4, 3, 5, 5)
        w = np.random.randn(2, 3, 3, 3)
        b = np.random.randn(2,)
        dout = np.random.randn(4, 2, 5, 5)
        conv_param = {'stride': 1, 'pad': 1}

        out, cache = conv_forward_naive(x,w,b,conv_param)

        dx_num = eval_numerical_gradient_array(lambda x:
        conv_forward_naive(x, w, b, conv_param)[0], x, dout)
        dw_num = eval_numerical_gradient_array(lambda w:
        conv_forward_naive(x, w, b, conv_param)[0], w, dout)
        db_num = eval_numerical_gradient_array(lambda b:
        conv_forward_naive(x, w, b, conv_param)[0], b, dout)

        out, cache = conv_forward_naive(x, w, b, conv_param)
        dx, dw, db = conv_backward_naive(dout, cache)

        # Your errors should be around 1e-9'
        print('Testing conv_backward_naive function')
        print('dx error: ', rel_error(dx, dx_num))
        print('dw error: ', rel_error(dw, dw_num))
        print('db error: ', rel_error(db, db_num))

Testing conv_backward_naive function
('dx error: ', 1.8325298354650547e-09)
('dw error: ', 2.1593755367307746e-09)
('db error: ', 2.832770724722921e-11)

```

Max pool forward pass

In this section, we will implement the forward pass of the max pool. The function is `max_pool_forward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_forward_naive`, test your implementation by running the cell below.

```
In [4]: x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)

correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                          [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]],
                          [[-0.02736842, -0.01263158],
                           [ 0.03157895,  0.04631579]]],
                        [[[ 0.09052632,  0.10526316],
                           [ 0.14947368,  0.16421053]],
                          [[ 0.20842105,  0.22315789],
                           [ 0.26736842,  0.28210526]],
                          [[ 0.32631579,  0.34105263],
                           [ 0.38526316,  0.4          ]]]]])

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))

Testing max_pool_forward_naive function:
('difference: ', 4.1666665157267834e-08)
```

Max pool backward pass

In this section, you will implement the backward pass of the max pool. The function is `max_pool_backward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_backward_naive`, test your implementation by running the cell below.

```
In [5]: x = np.random.randn(3, 2, 8, 8)
dout = np.random.randn(3, 2, 4, 4)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x, pool_param)[0], x, dout)

out, cache = max_pool_forward_naive(x, pool_param)
dx = max_pool_backward_naive(dout, cache)

# Your error should be around 1e-12
print('Testing max_pool_backward_naive function:')
print('dx error: ', rel_error(dx, dx_num))

Testing max_pool_backward_naive function:
('dx error: ', 3.2756139190913013e-12)
```

Fast implementation of the CNN layers

Implementing fast versions of the CNN layers can be difficult. We will provide you with the fast layers implemented by cs231n. They are provided in `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:

```
python setup.py build_ext --inplace
```

NOTE: The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the cell below.

You should see pretty drastic speedups in the implementation of these layers. On our machine, the forward pass speeds up by 17x and the backward pass speeds up by 840x. Of course, these numbers will vary from machine to machine, as well as on your precise implementation of the naive layers.

```

In [5]: from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
        from time import time

        x = np.random.randn(100, 3, 31, 31)
        w = np.random.randn(25, 3, 3, 3)
        b = np.random.randn(25,)
        dout = np.random.randn(100, 25, 16, 16)
        conv_param = {'stride': 2, 'pad': 1}

        t0 = time()
        out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
        t1 = time()
        out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
        t2 = time()

        print('Testing conv_forward_fast:')
        print('Naive: %fs' % (t1 - t0))
        print('Fast: %fs' % (t2 - t1))
        print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
        print('Difference: ', rel_error(out_naive, out_fast))

        t0 = time()
        dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
        t1 = time()
        dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
        t2 = time()

        print('\nTesting conv_backward_fast:')
        print('Naive: %fs' % (t1 - t0))
        print('Fast: %fs' % (t2 - t1))
        print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
        print('dx difference: ', rel_error(dx_naive, dx_fast))
        print('dw difference: ', rel_error(dw_naive, dw_fast))
        print('db difference: ', rel_error(db_naive, db_fast))

Testing conv_forward_fast:
Naive: 7.371448s
Fast: 0.023134s
Speedup: 318.641403x
('Difference: ', 9.798530848922286e-11)

Testing conv_backward_fast:
Naive: 13.187574s
Fast: 0.015983s
Speedup: 825.106956x
('dx difference: ', 1.4695323361077297e-11)
('dw difference: ', 4.5324381887239014e-13)
('db difference: ', 1.9028207917428035e-14)

```

```
In [6]: from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))

Testing pool_forward_fast:
Naive: 0.398780s
fast: 0.002665s
speedup: 149.647043x
('difference: ', 0.0)

Testing pool_backward_fast:
Naive: 0.642205s
speedup: 34.134263x
('dx difference: ', 0.0)
```

Implementation of cascaded layers

We've provided the following functions in `nndl/conv_layer_utils.py`:

- `conv_relu_forward`
- `conv_relu_backward`
- `conv_relu_pool_forward`
- `conv_relu_pool_backward`

These use the fast implementations of the conv net layers. You can test them below:

```
In [6]: from nndl.conv_layer_utils import conv_relu_pool_forward, conv_relu_p
ool_backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forwa
rd(x, w, b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forwa
rd(x, w, b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forwa
rd(x, w, b, conv_param, pool_param)[0], b, dout)

print('Testing conv_relu_pool')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

Testing conv_relu_pool
('dx error: ', 4.170875960314081e-08)
('dw error: ', 5.082348379616351e-09)
('db error: ', 4.4707648445387605e-10)
```



```
In [7]: from nndl.conv_layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x,
    w, b, conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x,
    w, b, conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x,
    w, b, conv_param)[0], b, dout)

print('Testing conv_relu:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

Testing conv_relu:
('dx error: ', 3.577807570434726e-09)
('dw error: ', 3.088580053034213e-10)
('db error: ', 1.5252153305647855e-10)
```

What next?

We saw how helpful batch normalization was for training FC nets. In the next notebook, we'll implement a batch normalization for convolutional neural networks, and then finish off by implementing a CNN to improve our validation accuracy on CIFAR-10.

Spatial batch normalization

In fully connected networks, we performed batch normalization on the activations. To do something equivalent on CNNs, we modify batch normalization slightly.

Normally batch-normalization accepts inputs of shape (N, D) and produces outputs of shape (N, D) , where we normalize across the minibatch dimension N . For data coming from convolutional layers, batch normalization accepts inputs of shape (N, C, H, W) and produces outputs of shape (N, C, H, W) where the N dimension gives the minibatch size and the (H, W) dimensions give the spatial size of the feature map.

How do we calculate the spatial averages? First, notice that for the C feature maps we have (i.e., the layer has C filters) that each of these ought to have its own batch norm statistics, since each feature map may be picking out very different features in the images. However, within a feature map, we may assume that across all inputs and across all locations in the feature map, there ought to be relatively similar first and second order statistics. Hence, one way to think of spatial batch-normalization is to reshape the (N, C, H, W) array as an $(N * H * W, C)$ array and perform batch normalization on this array.

Since spatial batch norm and batch normalization are similar, it'd be good to at this point also copy and paste our prior implemented layers from HW #4. Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- `layers.py` for your FC network layers, as well as `batchnorm` and `dropout`.
- `layer_utils.py` for your combined FC network layers.
- `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

If you use your prior implementations of the `batchnorm`, then your spatial `batchnorm` implementation may be very short. Our implementations of the forward and backward pass are each 6 lines of code.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

```
In [1]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Spatial batch normalization forward pass

Implement the forward pass, `spatial_batchnorm_forward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```
In [18]: # Check the training-time forward pass by checking means and variance
# of features both before and after spatial batch normalization

N, C, H, W = 2, 3, 4, 5
x = 4 * np.random.randn(N, C, H, W) + 10

print('Before spatial batch normalization:')
print(' Shape: ', x.shape)
print(' Means: ', x.mean(axis=(0, 2, 3)))
print(' Stds: ', x.std(axis=(0, 2, 3)))

# Means should be close to zero and stds close to one
gamma, beta = np.ones(C), np.zeros(C)
bn_param = {'mode': 'train'}
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization:')
print(' Shape: ', out.shape)
print(' Means: ', out.mean(axis=(0, 2, 3)))
print(' Stds: ', out.std(axis=(0, 2, 3)))

# Means should be close to beta and stds close to gamma
gamma, beta = np.asarray([3, 4, 5]), np.asarray([6, 7, 8])
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization (nontrivial gamma, beta):')
print(' Shape: ', out.shape)
print(' Means: ', out.mean(axis=(0, 2, 3)))
print(' Stds: ', out.std(axis=(0, 2, 3)))
```

```
Before spatial batch normalization:
(' Shape: ', (2, 3, 4, 5))
(' Means: ', array([10.63674602,  9.76709054, 12.15721725]))
(' Stds: ', array([4.19309582, 4.34852982, 4.5610162 ]))
After spatial batch normalization:
(' Shape: ', (2, 3, 4, 5))
(' Means: ', array([-2.28983499e-16,  4.57966998e-16, -3.62904151e-16]))
(' Stds: ', array([0.99999972, 0.99999974, 0.99999976]))
After spatial batch normalization (nontrivial gamma, beta):
(' Shape: ', (2, 3, 4, 5))
(' Means: ', array([6., 7., 8.]))
(' Stds: ', array([2.99999915, 3.99999894, 4.9999988 ]))
```

Spatial batch normalization backward pass

Implement the backward pass, `spatial_batchnorm_backward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```
In [19]: N, C, H, W = 2, 3, 4, 5
x = 5 * np.random.randn(N, C, H, W) + 12
gamma = np.random.randn(C)
beta = np.random.randn(C)
dout = np.random.randn(N, C, H, W)

bn_param = {'mode': 'train'}
fx = lambda x: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fg = lambda a: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fb = lambda b: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]

dx_num = eval_numerical_gradient_array(fx, x, dout)
da_num = eval_numerical_gradient_array(fg, gamma, dout)
db_num = eval_numerical_gradient_array(fb, beta, dout)

_, cache = spatial_batchnorm_forward(x, gamma, beta, bn_param)
dx, dgamma, dbeta = spatial_batchnorm_backward(dout, cache)
print('dx error: ', rel_error(dx_num, dx))
print('dgamma error: ', rel_error(da_num, dgamma))
print('dbeta error: ', rel_error(db_num, dbeta))

('dx error: ', 1.4042063030895136e-08)
('dgamma error: ', 2.4816141845255796e-12)
('dbeta error: ', 3.2775501014269754e-12)
```

Convolutional neural networks

In this notebook, we'll put together our convolutional layers to implement a 3-layer CNN. Then, we'll ask you to implement a CNN that can achieve $> 65\%$ validation error on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

If you have not completed the Spatial BatchNorm Notebook, please see the following description from that notebook:

Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their layer implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- `layers.py` for your FC network layers, as well as batchnorm and dropout.
- `layer_utils.py` for your combined FC network layers.
- `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

```

In [1]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from nndl.cnn import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient_array, eval_numerical_gradient
from nndl.layers import *
from nndl.conv_layers import *
from cs231n.fast_layers import *
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

```

In [2]: # Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))

X_val: (1000, 3, 32, 32)
X_train: (49000, 3, 32, 32)
X_test: (1000, 3, 32, 32)
y_val: (1000,)
y_train: (49000,)
y_test: (1000,)

```

Three layer CNN

In this notebook, you will implement a three layer CNN. The `ThreeLayerConvNet` class is in `nndl/cnn.py`. You'll need to modify that code for this section, including the initialization, as well as the calculation of the loss and gradients. You should be able to use the building blocks you have either earlier coded or that we have provided. Be sure to use the fast layers.

The architecture of this CNN will be:

conv - relu - 2x2 max pool - affine - relu - affine - softmax

We won't use batchnorm yet. You've also done enough of these to know how to debug; use the cells below.

Note: As we are implementing several layers CNN networks. The gradient error can be expected for the `eval_numerical_gradient()` function. If your W1 max relative error and W2 max relative error are around or below 0.01, they should be acceptable. Other errors should be less than 1e-5.

```
In [8]: num_inputs = 2
input_dim = (3, 16, 16)
reg = 0.0
num_classes = 10
X = np.random.randn(num_inputs, *input_dim)
y = np.random.randint(num_classes, size=num_inputs)

model = ThreeLayerConvNet(num_filters=3, filter_size=3,
                           input_dim=input_dim, hidden_dim=7,
                           dtype=np.float64)
loss, grads = model.loss(X, y)
for param_name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    param_grad_num = eval_numerical_gradient(f, model.params[param_name], verbose=False, h=1e-6)
    e = rel_error(param_grad_num, grads[param_name])
    print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_name])))

W1 max relative error: 0.00292987249321
W2 max relative error: 0.0155920132487
W3 max relative error: 7.11899484973e-05
b1 max relative error: 5.0406340691e-05
b2 max relative error: 1.92297083123e-05
b3 max relative error: 1.33771926975e-09
```

Overfit small dataset

To check your CNN implementation, let's overfit a small dataset.


```
In [9]: num_train = 100
        small_data = {
            'X_train': data['X_train'][:num_train],
            'y_train': data['y_train'][:num_train],
            'X_val': data['X_val'],
            'y_val': data['y_val'],
        }


        model = ThreeLayerConvNet(weight_scale=1e-2)

        solver = Solver(model, small_data,
                        num_epochs=20, batch_size=50,
                        update_rule='adam',
                        optim_config={
                            'learning_rate': 1e-3,
                        },
                        verbose=True, print_every=1)
        solver.train()
```

```
(Iteration 1 / 40) loss: 2.398724
(Epoch 0 / 20) train acc: 0.240000; val_acc: 0.123000
(Iteration 2 / 40) loss: 3.658866
(Epoch 1 / 20) train acc: 0.170000; val_acc: 0.113000
(Iteration 3 / 40) loss: 2.623158
(Iteration 4 / 40) loss: 2.581057
(Epoch 2 / 20) train acc: 0.230000; val_acc: 0.149000
(Iteration 5 / 40) loss: 2.198190
(Iteration 6 / 40) loss: 2.208519
(Epoch 3 / 20) train acc: 0.260000; val_acc: 0.119000
(Iteration 7 / 40) loss: 2.139135
(Iteration 8 / 40) loss: 1.961251
(Epoch 4 / 20) train acc: 0.370000; val_acc: 0.166000
(Iteration 9 / 40) loss: 1.886949
(Iteration 10 / 40) loss: 1.919783
(Epoch 5 / 20) train acc: 0.440000; val_acc: 0.165000
(Iteration 11 / 40) loss: 1.434302
(Iteration 12 / 40) loss: 1.668035
(Epoch 6 / 20) train acc: 0.540000; val_acc: 0.192000
(Iteration 13 / 40) loss: 1.352097
(Iteration 14 / 40) loss: 1.305671
(Epoch 7 / 20) train acc: 0.660000; val_acc: 0.206000
(Iteration 15 / 40) loss: 1.089647
(Iteration 16 / 40) loss: 1.030339
(Epoch 8 / 20) train acc: 0.720000; val_acc: 0.195000
(Iteration 17 / 40) loss: 0.933218
(Iteration 18 / 40) loss: 0.653353
(Epoch 9 / 20) train acc: 0.730000; val_acc: 0.195000
(Iteration 19 / 40) loss: 0.860156
(Iteration 20 / 40) loss: 0.956476
(Epoch 10 / 20) train acc: 0.820000; val_acc: 0.213000
(Iteration 21 / 40) loss: 0.673383
(Iteration 22 / 40) loss: 0.626792
(Epoch 11 / 20) train acc: 0.860000; val_acc: 0.203000
(Iteration 23 / 40) loss: 0.410349
(Iteration 24 / 40) loss: 0.297234
(Epoch 12 / 20) train acc: 0.890000; val_acc: 0.201000
(Iteration 25 / 40) loss: 0.280284
(Iteration 26 / 40) loss: 0.206168
(Epoch 13 / 20) train acc: 0.930000; val_acc: 0.180000
(Iteration 27 / 40) loss: 0.242705
(Iteration 28 / 40) loss: 0.332490
(Epoch 14 / 20) train acc: 0.970000; val_acc: 0.197000
(Iteration 29 / 40) loss: 0.138159
(Iteration 30 / 40) loss: 0.265789
(Epoch 15 / 20) train acc: 0.980000; val_acc: 0.211000
(Iteration 31 / 40) loss: 0.064454
(Iteration 32 / 40) loss: 0.102391
(Epoch 16 / 20) train acc: 0.990000; val_acc: 0.217000
(Iteration 33 / 40) loss: 0.049662
(Iteration 34 / 40) loss: 0.098386
(Epoch 17 / 20) train acc: 0.990000; val_acc: 0.217000
(Iteration 35 / 40) loss: 0.080395
(Iteration 36 / 40) loss: 0.069240
(Epoch 18 / 20) train acc: 1.000000; val_acc: 0.217000
(Iteration 37 / 40) loss: 0.026617
(Iteration 38 / 40) loss: 0.049096
```

```
(Epoch 19 / 20) train acc: 0.980000; val_acc: 0.223000  
(Iteration 39 / 40) loss: 0.038309  
(Iteration 40 / 40) loss: 0.021915  
(Epoch 20 / 20) train acc: 1.000000; val_acc: 0.215000
```

```
In [10]: plt.subplot(2, 1, 1)  
plt.plot(solver.loss_history, 'o')  
plt.xlabel('iteration')  
plt.ylabel('loss')  
  
plt.subplot(2, 1, 2)  
plt.plot(solver.train_acc_history, '-o')  
plt.plot(solver.val_acc_history, '-o')  
plt.legend(['train', 'val'], loc='upper left')  
plt.xlabel('epoch')  
plt.ylabel('accuracy')  
plt.show()
```



Train the network

Now we train the 3 layer CNN on CIFAR-10 and assess its accuracy.

```
In [11]: model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=1, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)
solver.train()
```

```
(Iteration 1 / 980) loss: 2.304511
(Epoch 0 / 1) train acc: 0.128000; val_acc: 0.139000
(Iteration 21 / 980) loss: 2.231854
(Iteration 41 / 980) loss: 2.148167
(Iteration 61 / 980) loss: 1.939863
(Iteration 81 / 980) loss: 1.755125
(Iteration 101 / 980) loss: 1.797583
(Iteration 121 / 980) loss: 1.717661
(Iteration 141 / 980) loss: 1.942601
(Iteration 161 / 980) loss: 1.737939
(Iteration 181 / 980) loss: 1.765807
(Iteration 201 / 980) loss: 1.970439
(Iteration 221 / 980) loss: 1.943240
(Iteration 241 / 980) loss: 1.789837
(Iteration 261 / 980) loss: 1.686366
(Iteration 281 / 980) loss: 1.725475
(Iteration 301 / 980) loss: 1.504726
(Iteration 321 / 980) loss: 1.849093
(Iteration 341 / 980) loss: 1.331362
(Iteration 361 / 980) loss: 2.033096
(Iteration 381 / 980) loss: 1.936199
(Iteration 401 / 980) loss: 1.587918
(Iteration 421 / 980) loss: 1.521933
(Iteration 441 / 980) loss: 1.531843
(Iteration 461 / 980) loss: 1.640172
(Iteration 481 / 980) loss: 1.494496
(Iteration 501 / 980) loss: 1.694095
(Iteration 521 / 980) loss: 1.676323
(Iteration 541 / 980) loss: 1.590233
(Iteration 561 / 980) loss: 1.577411
(Iteration 581 / 980) loss: 1.630510
(Iteration 601 / 980) loss: 1.802195
(Iteration 621 / 980) loss: 1.790504
(Iteration 641 / 980) loss: 1.791112
(Iteration 661 / 980) loss: 1.720310
(Iteration 681 / 980) loss: 1.328894
(Iteration 701 / 980) loss: 1.714998
(Iteration 721 / 980) loss: 1.439794
(Iteration 741 / 980) loss: 1.322611
(Iteration 761 / 980) loss: 1.479585
(Iteration 781 / 980) loss: 1.402265
(Iteration 801 / 980) loss: 1.576688
(Iteration 821 / 980) loss: 1.518877
(Iteration 841 / 980) loss: 1.505881
(Iteration 861 / 980) loss: 1.708602
(Iteration 881 / 980) loss: 1.465694
(Iteration 901 / 980) loss: 1.518513
(Iteration 921 / 980) loss: 1.471163
(Iteration 941 / 980) loss: 1.342395
(Iteration 961 / 980) loss: 1.610166
(Epoch 1 / 1) train acc: 0.494000; val_acc: 0.472000
```

Get > 65% validation accuracy on CIFAR-10.

In the last part of the assignment, we'll now ask you to train a CNN to get better than 65% validation accuracy on CIFAR-10.

Things you should try:

- Filter size: Above we used 7x7; but VGGNet and onwards showed stacks of 3x3 filters are good.
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Batch normalization: Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- Network architecture: Can a deeper CNN do better? Consider these architectures:
 - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
 - [conv-relu-pool]XN - [affine]XM - [softmax or SVM]
 - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

```
In [3]: # ===== #
# YOUR CODE HERE:
#   Implement a CNN to achieve greater than 65% validation accuracy
#   on CIFAR-10.
# ===== #
model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001,
                           num_filters=32)

solver = Solver(model, data,
                 num_epochs=20, batch_size=200,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)

solver.train()

# ===== #
# END YOUR CODE HERE
# ===== #
```

```
(Iteration 1 / 4900) loss: 2.304641
(Epoch 0 / 20) train acc: 0.099000; val_acc: 0.087000
(Iteration 21 / 4900) loss: 2.022833
(Iteration 41 / 4900) loss: 1.829609
(Iteration 61 / 4900) loss: 1.578786
(Iteration 81 / 4900) loss: 1.525048
(Iteration 101 / 4900) loss: 1.524534
(Iteration 121 / 4900) loss: 1.491584
(Iteration 141 / 4900) loss: 1.484886
(Iteration 161 / 4900) loss: 1.451082
(Iteration 181 / 4900) loss: 1.384444
(Iteration 201 / 4900) loss: 1.391388
(Iteration 221 / 4900) loss: 1.263857
(Iteration 241 / 4900) loss: 1.477539
(Epoch 1 / 20) train acc: 0.507000; val_acc: 0.493000
(Iteration 261 / 4900) loss: 1.262600
(Iteration 281 / 4900) loss: 1.369997
(Iteration 301 / 4900) loss: 1.176921
(Iteration 321 / 4900) loss: 1.533619
(Iteration 341 / 4900) loss: 1.322418
(Iteration 361 / 4900) loss: 1.153445
(Iteration 381 / 4900) loss: 1.329984
(Iteration 401 / 4900) loss: 1.336941
(Iteration 421 / 4900) loss: 1.231730
(Iteration 441 / 4900) loss: 1.078800
(Iteration 461 / 4900) loss: 1.286022
(Iteration 481 / 4900) loss: 1.337707
(Epoch 2 / 20) train acc: 0.605000; val_acc: 0.584000
(Iteration 501 / 4900) loss: 1.134415
(Iteration 521 / 4900) loss: 1.144201
(Iteration 541 / 4900) loss: 1.010227
(Iteration 561 / 4900) loss: 1.293241
(Iteration 581 / 4900) loss: 1.193454
(Iteration 601 / 4900) loss: 1.102433
(Iteration 621 / 4900) loss: 1.153127
(Iteration 641 / 4900) loss: 0.996380
(Iteration 661 / 4900) loss: 1.083883
(Iteration 681 / 4900) loss: 1.194233
(Iteration 701 / 4900) loss: 1.146079
(Iteration 721 / 4900) loss: 1.307219
(Epoch 3 / 20) train acc: 0.626000; val_acc: 0.603000
(Iteration 741 / 4900) loss: 1.176179
(Iteration 761 / 4900) loss: 1.221810
(Iteration 781 / 4900) loss: 1.322550
(Iteration 801 / 4900) loss: 1.156965
(Iteration 821 / 4900) loss: 1.254104
(Iteration 841 / 4900) loss: 0.918489
(Iteration 861 / 4900) loss: 1.142895
(Iteration 881 / 4900) loss: 1.190759
(Iteration 901 / 4900) loss: 1.099620
(Iteration 921 / 4900) loss: 1.129262
(Iteration 941 / 4900) loss: 1.066536
(Iteration 961 / 4900) loss: 1.286184
(Epoch 4 / 20) train acc: 0.636000; val_acc: 0.624000
(Iteration 981 / 4900) loss: 1.030652
(Iteration 1001 / 4900) loss: 1.120823
(Iteration 1021 / 4900) loss: 1.030835
```



```
(Iteration 1041 / 4900) loss: 1.124149
(Iteration 1061 / 4900) loss: 1.102276
(Iteration 1081 / 4900) loss: 0.881765
(Iteration 1101 / 4900) loss: 1.156674
(Iteration 1121 / 4900) loss: 1.233775
(Iteration 1141 / 4900) loss: 1.089831
(Iteration 1161 / 4900) loss: 0.990193
(Iteration 1181 / 4900) loss: 1.056324
(Iteration 1201 / 4900) loss: 1.150618
(Iteration 1221 / 4900) loss: 0.922874
(Epoch 5 / 20) train acc: 0.708000; val_acc: 0.631000
(Iteration 1241 / 4900) loss: 0.953913
(Iteration 1261 / 4900) loss: 0.953533
(Iteration 1281 / 4900) loss: 1.131170
(Iteration 1301 / 4900) loss: 1.108464
(Iteration 1321 / 4900) loss: 0.915440
(Iteration 1341 / 4900) loss: 0.964860
(Iteration 1361 / 4900) loss: 1.069142
(Iteration 1381 / 4900) loss: 1.014813
(Iteration 1401 / 4900) loss: 1.015884
(Iteration 1421 / 4900) loss: 1.091674
(Iteration 1441 / 4900) loss: 0.916118
(Iteration 1461 / 4900) loss: 0.929332
(Epoch 6 / 20) train acc: 0.691000; val_acc: 0.624000
(Iteration 1481 / 4900) loss: 0.874445
(Iteration 1501 / 4900) loss: 0.853501
(Iteration 1521 / 4900) loss: 0.960778
(Iteration 1541 / 4900) loss: 0.990698
(Iteration 1561 / 4900) loss: 1.192997
(Iteration 1581 / 4900) loss: 1.125338
(Iteration 1601 / 4900) loss: 0.965832
(Iteration 1621 / 4900) loss: 1.030510
(Iteration 1641 / 4900) loss: 1.083969
(Iteration 1661 / 4900) loss: 0.859961
(Iteration 1681 / 4900) loss: 0.886632
(Iteration 1701 / 4900) loss: 0.962145
(Epoch 7 / 20) train acc: 0.691000; val_acc: 0.597000
(Iteration 1721 / 4900) loss: 0.983099
(Iteration 1741 / 4900) loss: 0.940587
(Iteration 1761 / 4900) loss: 1.006831
(Iteration 1781 / 4900) loss: 0.881595
(Iteration 1801 / 4900) loss: 0.949778
(Iteration 1821 / 4900) loss: 0.872129
(Iteration 1841 / 4900) loss: 0.973894
(Iteration 1861 / 4900) loss: 0.881316
(Iteration 1881 / 4900) loss: 0.952845
(Iteration 1901 / 4900) loss: 0.882571
(Iteration 1921 / 4900) loss: 0.979896
(Iteration 1941 / 4900) loss: 0.956969
(Epoch 8 / 20) train acc: 0.717000; val_acc: 0.652000
(Iteration 1961 / 4900) loss: 0.892523
(Iteration 1981 / 4900) loss: 0.834394
(Iteration 2001 / 4900) loss: 0.917589
```

```

-----
-----
KeyboardInterrupt                                Traceback (most recent call
last)
<ipython-input-3-1e80315e6404> in <module>()
    15             verbose=True, print_every=20)
    16
--> 17 solver.train()
    18
    19 # =====
===== #

/home/ben/Documents/239AS/HW5/code/cs231n/solver.pyc in train(self)
    262
    263         for t in range(num_iterations):
--> 264             self._step()
    265
    266             # Maybe print training loss

/home/ben/Documents/239AS/HW5/code/cs231n/solver.pyc in _step(self)
    178
    179         # Compute loss and gradient
--> 180         loss, grads = self.model.loss(X_batch, y_batch)
    181         self.loss_history.append(loss)
    182

/home/ben/Documents/239AS/HW5/code/nndl/cnn.pyc in loss(self, X, y)
    102     # =====
===== #
    103
--> 104     c1, conv_cache1 = conv_forward_fast(X, W1, b1,
conv_param)
    105     h1, r_cache1 = relu_forward(c1)
    106     mp1, mp_cache1 = max_pool_forward_fast(h1, pool_param)

/home/ben/Documents/239AS/HW5/code/cs231n/fast_layers.pyc in conv_for
ward_strides(x, w, b, conv_param)
    72
    73     # Now all our convolutions are a big matrix multiply
--> 74     res = w.reshape(F, -1).dot(x_cols) + b.reshape(-1, 1)
    75
    76     # Reshape the output

```

KeyboardInterrupt:

In [4]: **print** solver.best_val_acc

0.652

```

1 import numpy as np
2 from nndl.layers import *
3 import pdb
4
5 """
6 This code was originally written for CS 231n at Stanford University
7 (cs231n.stanford.edu). It has been modified in various areas for use in the
8 ECE 239AS class at UCLA. This includes the descriptions of what code to
9 implement as well as some slight potential changes in variable names to be
10 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 def conv_forward_naive(x, w, b, conv_param):
16     """
17     A naive implementation of the forward pass for a convolutional layer.
18
19     The input consists of N data points, each with C channels, height H and width
20     W. We convolve each input with F different filters, where each filter spans
21     all C channels and has height HH and width WW.
22
23     Input:
24     - x: Input data of shape (N, C, H, W)
25     - w: Filter weights of shape (F, C, HH, WW)
26     - b: Biases, of shape (F,)
27     - conv_param: A dictionary with the following keys:
28       - 'stride': The number of pixels between adjacent receptive fields in the
29         horizontal and vertical directions.
30       - 'pad': The number of pixels that will be used to zero-pad the input.
31
32     Returns a tuple of:
33     - out: Output data, of shape (N, F, H', W') where H' and W' are given by
34        $H' = 1 + (H + 2 * \text{pad} - \text{HH}) / \text{stride}$ 
35        $W' = 1 + (W + 2 * \text{pad} - \text{WW}) / \text{stride}$ 
36     - cache: (x, w, b, conv_param)
37     """
38     out = None
39     pad = conv_param['pad']
40     stride = conv_param['stride']
41
42     # ===== #
43     # YOUR CODE HERE:
44     # Implement the forward pass of a convolutional neural network.
45     # Store the output as 'out'.
46     # Hint: to pad the array, you can use the function np.pad.
47     # ===== #
48
49     N, C, H, W = x.shape
50     F, _, HH, WW = w.shape
51     npad = ((0,0), (0,0), (pad,pad), (pad,pad))
52     x_padded = np.pad(x, pad_width=npad, mode='constant', constant_values=0)
53     H_prime = 1 + (H + 2*pad - HH)/stride
54     W_prime = 1 + (W + 2*pad - WW)/stride
55
56     out = np.empty((N,F,H_prime, W_prime))
57
58     for n in np.arange(0,N):
59         for i in np.arange(0, F):
60             for j in np.arange(0, H_prime):
61                 for k in np.arange(0, W_prime):
62                     out[n,i,j,k] = np.sum(w[i, :, :, :] * x_padded[n, :, j*stride:j*stride+HH, k*stride:k*stride+WW]) + b[i]
63
64
65     # ===== #
66     # END YOUR CODE HERE
67     # ===== #
68
69     cache = (x, w, b, conv_param)
70     return out, cache
71
72
73
74 def conv_backward_naive(dout, cache):
75     """
76     A naive implementation of the backward pass for a convolutional layer.
77
78     Inputs:
79     - dout: Upstream derivatives.
80     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
81

```

```

82     Returns a tuple of:
83     - dx: Gradient with respect to x
84     - dw: Gradient with respect to w
85     - db: Gradient with respect to b
86     """
87     dx, dw, db = None, None, None
88
89     N, F, out_height, out_width = dout.shape
90     x, w, b, conv_param = cache
91
92     stride, pad = [conv_param['stride'], conv_param['pad']]
93     xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
94     num_filts, _, f_height, f_width = w.shape
95
96     # ===== #
97     # YOUR CODE HERE:
98     # Implement the backward pass of a convolutional neural network.
99     # Calculate the gradients: dx, dw, and db.
100    # ===== #
101
102    dw = np.zeros_like(w)
103    db = np.zeros_like(b)
104    dx = np.zeros_like(x)
105    dxpad = np.zeros_like(xpad)
106    F, _, HH, WW = w.shape
107
108    for n in np.arange(0,N):
109        for f in np.arange(0, F):
110            for j in np.arange(0, out_height):
111                for k in np.arange(0, out_width):
112                    dw[f] += xpad[n, :, j*stride:j*stride+HH, k*stride:k*stride+WW]*dout[n, f, j, k]
113                    db[f] += dout[n,f,j,k]
114                    dxpad[n, :, j*stride:j*stride+HH, k*stride:k*stride+WW] += w[f]*dout[n,f,j,k]
115
116    dx[:] = dxpad[:, :, pad:-pad, pad:-pad]
117
118
119
120
121
122    # ===== #
123    # END YOUR CODE HERE
124    # ===== #
125
126    return dx, dw, db
127
128
129 def max_pool_forward_naive(x, pool_param):
130     """
131     A naive implementation of the forward pass for a max pooling layer.
132
133     Inputs:
134     - x: Input data, of shape (N, C, H, W)
135     - pool_param: dictionary with the following keys:
136       - 'pool_height': The height of each pooling region
137       - 'pool_width': The width of each pooling region
138       - 'stride': The distance between adjacent pooling regions
139
140     Returns a tuple of:
141     - out: Output data
142     - cache: (x, pool_param)
143     """
144     out = None
145
146     # ===== #
147     # YOUR CODE HERE:
148     # Implement the max pooling forward pass.
149     # ===== #
150     N, C, H, W = x.shape
151
152     ph = pool_param['pool_height']
153     pw = pool_param['pool_width']
154     stride = pool_param['stride']
155     h_prime = (H - ph)/stride + 1
156     w_prime = (W - pw)/stride + 1
157
158     out = np.empty((N,C,h_prime, w_prime))
159
160     for n in np.arange(0,N):
161         for c in np.arange(0,C):
162             for h in np.arange(0,h_prime):
163                 for w in np.arange(0,w_prime):

```

```

164         out[n,c,h,w] = np.max(x[n,c,h*stride:h*stride+ph,w*stride:w*stride+pw])
165
166
167         # ===== #
168         # END YOUR CODE HERE
169         # ===== #
170     cache = (x, pool_param)
171     return out, cache
172
173 def max_pool_backward_naive(dout, cache):
174     """
175     A naive implementation of the backward pass for a max pooling layer.
176
177     Inputs:
178     - dout: Upstream derivatives
179     - cache: A tuple of (x, pool_param) as in the forward pass.
180
181     Returns:
182     - dx: Gradient with respect to x
183     """
184     dx = None
185     x, pool_param = cache
186     pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'], pool_param['stride']
187
188     # ===== #
189     # YOUR CODE HERE:
190     # Implement the max pooling backward pass.
191     # ===== #
192     N, C, H, W = x.shape
193     h_prime, w_prime = dout.shape[-2:]
194     dx = np.zeros_like(x)
195
196     for n in np.arange(0,N):
197         for c in np.arange(0,C):
198             for h in np.arange(0,h_prime):
199                 for w in np.arange(0,w_prime):
200                     idx = np.unravel_index( np.argmax( x[n, c, h*stride:h*stride+pool_height, w*stride:w*stride+pool_width]),
201                     (pool_height, pool_width))
202                     dx[n,c,h*stride+idx[0],w*stride+idx[1]] = dout[n,c,h,w]
203
204
205
206     # ===== #
207     # END YOUR CODE HERE
208     # ===== #
209
210     return dx
211
212 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
213     """
214     Computes the forward pass for spatial batch normalization.
215
216     Inputs:
217     - x: Input data of shape (N, C, H, W)
218     - gamma: Scale parameter, of shape (C,)
219     - beta: Shift parameter, of shape (C,)
220     - bn_param: Dictionary with the following keys:
221       - mode: 'train' or 'test'; required
222       - eps: Constant for numeric stability
223       - momentum: Constant for running mean / variance. momentum=0 means that
224         old information is discarded completely at every time step, while
225         momentum=1 means that new information is never incorporated. The
226         default of momentum=0.9 should work well in most situations.
227       - running_mean: Array of shape (D,) giving running mean of features
228       - running_var: Array of shape (D,) giving running variance of features
229
230     Returns a tuple of:
231     - out: Output data, of shape (N, C, H, W)
232     - cache: Values needed for the backward pass
233     """
234     out, cache = None, None
235
236     # ===== #
237     # YOUR CODE HERE:
238     # Implement the spatial batchnorm forward pass.
239     #
240     # You may find it useful to use the batchnorm forward pass you
241     # implemented in HW #4.
242     # ===== #
243
244     x_transpose = x.transpose((0,2,3,1))
245     x_reshaped = x_transpose.reshape((-1,x.shape[1]))

```

```

246
247 out, cache = batchnorm_forward(x_resaped, gamma, beta, bn_param)
248
249 out = out.reshape(*x_transpose.shape).transpose((0,3,1,2))
250
251 # ===== #
252 # END YOUR CODE HERE
253 # ===== #
254
255 return out, cache
256
257
258 def spatial_batchnorm_backward(dout, cache):
259     """
260     Computes the backward pass for spatial batch normalization.
261
262     Inputs:
263     - dout: Upstream derivatives, of shape (N, C, H, W)
264     - cache: Values from the forward pass
265
266     Returns a tuple of:
267     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
268     - dgamma: Gradient with respect to scale parameter, of shape (C,)
269     - dbeta: Gradient with respect to shift parameter, of shape (C,)
270     """
271     dx, dgamma, dbeta = None, None, None
272
273     # ===== #
274     # YOUR CODE HERE:
275     # Implement the spatial batchnorm backward pass.
276     #
277     # You may find it useful to use the batchnorm forward pass you
278     # implemented in HW #4.
279     # ===== #
280
281     dout_t = dout.transpose(0,2,3,1)
282     dout_r = dout_t.reshape(-1, dout.shape[1])
283
284     dx, dgamma, dbeta = batchnorm_backward(dout_r, cache)
285
286     dx = dx.reshape(*dout_t.shape).transpose((0,3,1,2))
287
288     # ===== #
289     # END YOUR CODE HERE
290     # ===== #
291
292     return dx, dgamma, dbeta

```

```

1 import numpy as np
2
3 from nndl.layers import *
4 from nndl.conv_layers import *
5 from cs231n.fast_layers import *
6 from nndl.layer_utils import *
7 from nndl.conv_layer_utils import *
8
9 import pdb
10
11 """
12 This code was originally written for CS 231n at Stanford University
13 (cs231n.stanford.edu). It has been modified in various areas for use in the
14 ECE 239AS class at UCLA. This includes the descriptions of what code to
15 implement as well as some slight potential changes in variable names to be
16 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
17 permission to use this code. To see the original version, please visit
18 cs231n.stanford.edu.
19 """
20
21 class ThreeLayerConvNet(object):
22     """
23     A three-layer convolutional network with the following architecture:
24
25     conv - relu - 2x2 max pool - affine - relu - affine - softmax
26
27     The network operates on minibatches of data that have shape (N, C, H, W)
28     consisting of N images, each with height H and width W and with C input
29     channels.
30     """
31
32     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
33                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
34                 dtype=np.float32, use_batchnorm=False):
35         """
36         Initialize a new network.
37
38         Inputs:
39         - input_dim: Tuple (C, H, W) giving size of input data
40         - num_filters: Number of filters to use in the convolutional layer
41         - filter_size: Size of filters to use in the convolutional layer
42         - hidden_dim: Number of units to use in the fully-connected hidden layer
43         - num_classes: Number of scores to produce from the final affine layer.
44         - weight_scale: Scalar giving standard deviation for random initialization
45           of weights.
46         - reg: Scalar giving L2 regularization strength
47         - dtype: numpy datatype to use for computation.
48         """
49         self.use_batchnorm = use_batchnorm
50         self.params = {}
51         self.reg = reg
52         self.dtype = dtype
53
54         # ===== #
55         # YOUR CODE HERE:
56         # Initialize the weights and biases of a three layer CNN. To initialize:
57         # - the biases should be initialized to zeros.
58         # - the weights should be initialized to a matrix with entries
59         #   drawn from a Gaussian distribution with zero mean and
60         #   standard deviation given by weight_scale.
61         # ===== #
62
63         self.params['W1'] = weight_scale * np.random.randn(num_filters, input_dim[0], filter_size, filter_size)
64         self.params['b1'] = np.zeros(num_filters)
65         self.params['W2'] = weight_scale * np.random.randn(num_filters*(input_dim[1]/2)*(input_dim[2]/2), hidden_dim)
66         self.params['b2'] = np.zeros(hidden_dim)
67         self.params['W3'] = weight_scale * np.random.randn(hidden_dim, num_classes)
68         self.params['b3'] = np.zeros(num_classes)
69
70         # ===== #
71         # END YOUR CODE HERE
72         # ===== #
73
74         for k, v in self.params.items():
75             self.params[k] = v.astype(dtype)
76
77
78
79     def loss(self, X, y=None):

```

```

80     """
81     Evaluate loss and gradient for the three-layer convolutional network.
82
83     Input / output: Same API as TwoLayerNet in fc_net.py.
84     """
85     W1, b1 = self.params['W1'], self.params['b1']
86     W2, b2 = self.params['W2'], self.params['b2']
87     W3, b3 = self.params['W3'], self.params['b3']
88
89     # pass conv_param to the forward pass for the convolutional layer
90     filter_size = W1.shape[2]
91     conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
92
93     # pass pool_param to the forward pass for the max-pooling layer
94     pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
95
96     scores = None
97
98     # ===== #
99     # YOUR CODE HERE:
100    # Implement the forward pass of the three layer CNN. Store the output
101    # scores as the variable "scores".
102    # ===== #
103
104    c1, conv_cache1 = conv_forward_fast(X, W1, b1, conv_param)
105    h1, r_cache1 = relu_forward(c1)
106    mp1, mp_cache1 = max_pool_forward_fast(h1, pool_param)
107    h2, aff_cache1 = affine_relu_forward(mp1, W2, b2)
108    scores, aff_cache2 = affine_forward(h2, W3, b3)
109
110    # ===== #
111    # END YOUR CODE HERE
112    # ===== #
113
114    if y is None:
115        return scores
116
117    loss, grads = 0, {}
118    # ===== #
119    # YOUR CODE HERE:
120    # Implement the backward pass of the three layer CNN. Store the grads
121    # in the grads dictionary, exactly as before (i.e., the gradient of
122    # self.params[k] will be grads[k]). Store the loss as "loss", and
123    # don't forget to add regularization on ALL weight matrices.
124    # ===== #
125
126    loss, dout = softmax_loss(scores, y)
127    loss += .5*self.reg*np.sum(W1*W1) + .5*self.reg*np.sum(W2*W2) + .5*self.reg*np.sum(W3*W3)
128
129    dx3, dw3, db3 = affine_backward(dout, aff_cache2)
130    dx2, dw2, db2 = affine_relu_backward(dx3, aff_cache1)
131    dmp = max_pool_backward_fast(dx2, mp_cache1)
132    dr = relu_backward(dmp, r_cache1)
133    dx1, dw1, db1 = conv_backward_fast(dr, conv_cache1)
134
135
136    grads['W1'] = dw1 + self.reg*W1
137    grads['b1'] = db1
138    grads['W2'] = dw2 + self.reg*W2
139    grads['b2'] = db2
140    grads['W3'] = dw3 + self.reg*W3
141    grads['b3'] = db3
142
143
144    # ===== #
145    # END YOUR CODE HERE
146    # ===== #
147
148    return loss, grads
149
150

```